

Real-Time Event-Driven Data Pipeline for an E-Commerce shop

Project Brief

You are tasked with designing and building a **real-time, event-driven data pipeline** to support operational analytics for an e-commerce platform.

The company deals with continuous inflow of transactional data (orders, products, etc.) in flat file format. These files arrive randomly in an **Amazon S3** bucket. The goal is to immediately validate, transform, and compute business KPIs, and store results for real-time querying via **Amazon DynamoDB**.

This is a **production-style project** requiring clear architecture, clean data flow, and automated orchestration. You must use AWS-native services in a containerized setup.

Core Services You Must Use

Your solution **must** use all four of the following services:

| Service | Purpose |
|--------------------|--|
| Amazon S3 | Input file storage |
| Amazon ECS | Containerized data validation & transformation |
| AWS Step Functions | Workflow orchestration |
| Amazon DynamoDB | Real-time storage of KPIs and metrics |

You are free to include other AWS services **as needed**

Your Mission

Build a data pipeline that:

1. **Detects** new data files arriving in S3.
2. **Validates** the contents using a containerized service.
3. **Transforms** valid data into business KPIs.
4. **Stores** KPIs in DynamoDB using an optimized schema.
5. **Archives** or logs processed files and outcomes.
6. **Automates** the entire workflow using Step Functions.

You are expected to make **architectural decisions** and simulate **realistic use cases**.

Required KPIs

You must compute and store the following two types of KPIs in DynamoDB as part of your transformation logic.

1. Category-Level KPIs (Per Category, Per Day)

| Field | Description |
|-----------------|--|
| category | Product category (e.g., Electronics) |
| order_date | Date of the summarized orders |
| daily_revenue | Total revenue from that category for the day |
| avg_order_value | Average value of individual orders in the category |
| avg_return_rate | Percentage of returned orders for the category |

Store in a **DynamoDB table** optimized for querying by category and order_date.

2. Order-Level KPIs (Per Day)

| Field | Description |
|------------------|--|
| order_date | Date of the summarized orders |
| total_orders | Count of unique orders |
| total_revenue | Total revenue from all orders |
| total_items_sold | Total number of items sold |
| return_rate | Percentage of orders that were returned |
| unique_customers | Number of distinct customers who placed orders |

Store in a **separate DynamoDB table** optimized for querying by order_date.

Project Expectations

You are expected to:

- Demonstrate a deep understanding of **event-driven design**.
- Build a solution that can scale with **realistic data volumes**.
- Handle **validation, transformation, and storage** cleanly and modularly.
- Explain and justify all design choices.

System Requirements

ECS Tasks

- **Validation Task:**
 - Reject malformed or incomplete data.
 - Exit the pipeline gracefully on failure.
- **Transformation Task:**
 - Process clean data and compute KPIs.
 - Store results in DynamoDB.

Step Functions

- Orchestrate ECS tasks.
- Include failure paths, branching logic, and timeouts.

DynamoDB

- Design a table (or tables) optimized for querying KPIs.
- Use proper partition keys, sort keys, and secondary indexes if needed.

Logging & Monitoring

- Use **CloudWatch Logs** to track ECS task execution.
- Create error logs or alerts as appropriate.

Optional Bonus Challenges

These are recommended, and will push your learning even further:

- Add a **notification system** (e.g., email alert) for failures.
- Implement a **retry mechanism** or **dead-letter queue** for failed tasks.
- Add a simple **dashboard** to expose KPIs (can be mocked).

Need Help?

This project is designed to challenge you. If you're stuck:

- Start by sketching your **architecture on paper**.
- Break it down into services: **S3 → ECS → Step Functions → DynamoDB**.
- Simulate manually before fully automating.
- Reach out with specific blockers—I'll guide, not hand-hold.

I'm not sure where this should be; for the sake of consistency, please make sure your documentation includes **AT LEAST** the following points

- Your data format and sample schema
- Validation rules (e.g., missing fields, referential integrity)
- DynamoDB table structure and access patterns
- Step Function workflow explanation
- Any error-handling, retry, or logging logic
- Instructions to simulate or test the pipeline manually